

# Final Report — Loan Policy Project

## 1. Overview

This report summarizes the EDA, supervised predictive model, and offline RL experiments.

## 2. Supervised DL model (MLP)

- Best MLP config: {'hidden\_dims': [512, 256], 'dropout': 0.2, 'lr': 0.001, 'batch\_size': 512, 'epochs': 10, 'patience': 3}
- Test AUC: 0.7367
- Best Test F1: 0.4539 at threshold 0.26

## 3. Offline RL agent (Discrete CQL)

- Estimated policy value (mean reward per case on test): -0.15966403481283487

## 4. Why these metrics?

### AUC & F1 for DL model:

- AUC measures the model discrimination across thresholds, useful for ranking risk.
- F1 measures balance between precision and recall at a chosen threshold; important when class imbalance exists and when we care about correct default detection.

### Estimated Policy Value for RL agent:

- In offline RL we aim to learn a policy that maximizes expected return. Estimated policy value (average monetary reward on a held-out test set) directly measures the downstream business objective (profit minus losses).

## 5. Policy comparison and example differences

The supervised DL model defines a thresholding policy (approve when predicted default probability  $< t$ ). The RL agent learns a policy that maximizes expected reward and may approve some high-risk applicants if the expected reward (interest) outweighs expected losses. Examples with differing decisions can be inspected by comparing model predictions and RL action for the same applicant.

## 6. Limitations & Future Steps

- Behavior policy was simulated because raw logged approve/deny decisions were unavailable. Replace with real logged propensities if available.
- Use OPE (IS/DR) for robust policy evaluation, and tune CQL/IQL for better performance.
- Consider reward normalization (e.g., divide monetary values by loan amount) and more complex reward that includes time-discounted cashflows.

## 7. Reproducibility

Follow the instructions in 'README.md'. Key scripts:

- 'train\_mlp\_grid.py' — grid search for MLPs
- 'evaluate\_best.py' — evaluate and save final supervised metrics
- 'build\_rl\_dataset.py' and 'train\_offline\_rl.py' — offline RL pipeline